

Mental Health and Social Media Balance

1. Project Overview

This dataset captures the delicate relationship between social media habits and mental well-being. It combines variables such as screen time, stress level, sleep quality, digital detox days, and happiness index. Ideal for regression, correlation, or mental health prediction tasks.

2. Dataset Summary

Link - [Mental Health and Social Media Balance](#)

This dataset is from Kaggle. It contains Mental Health report and social media Balance. Researchers, psychologists, and data enthusiasts can use this dataset to study how lifestyle and online activity patterns affect human emotions and overall wellness

- Rows 500
- Columns 10

- Key features:

| | |
|--------------------|-----------------------------------|
| User ID | - Unique identifier for each user |
| Age | - Age of users |
| Gender | - Gender of Users |
| Daily Screen Time | - Daily Scree time of users |
| Sleep Quality | - Sleep Quality of Users |
| Stress Level | - Geographic sales region |
| Days Without Media | - Types of stores |
| Exercise | - Exercise habits of users |
| Social Media | - Social media usage of users |
| Happiness Index | - Total happiness index out of 10 |

3. Exploratory Data Analysis using Python

We began with data preparation and cleaning in python

*Defining the columns

```
Index(['User_ID', 'Age', 'Gender', 'Daily_Screen_Time(hrs)',  
       'Sleep_Quality(1-10)', 'Stress_Level(1-10)',  
       'Days_Without_Social_Media', 'Exercise_Frequency(week)',  
       'Social_Media_Platform', 'Happiness_Index(1-10)'),  
      dtypes='object')
```

*Initial Exploration: Used df.info () to check structure.

```
Data columns (total 10 columns):  
 #   Column           Non-Null Count  Dtype     
 ---    
 0   User_ID          500 non-null    object    
 1   Age              500 non-null    int64    
 2   Gender            500 non-null    object    
 3   Daily_Screen_Time(hrs) 500 non-null  float64  
 4   Sleep_Quality(1-10) 500 non-null  float64  
 5   Stress_Level(1-10) 500 non-null  float64  
 6   Days_Without_Social_Media 500 non-null  float64  
 7   Exercise_Frequency(week) 500 non-null  float64  
 8   Social_Media_Platform 500 non-null  object    
 9   Happiness_Index(1-10) 500 non-null  float64  
 dtypes: float64(6), int64(1), object(3)
```

* Describe () for summary statistics.

| | Age | Daily_Screen_Time(hrs) | Sleep_Quality(1-10) | Stress_Level(1-10) | Days_Without_Social_Media | Exercise_Frequency(week) |
|-------|------------|------------------------|---------------------|--------------------|---------------------------|--------------------------|
| count | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000 |
| mean | 32.988000 | 5.530000 | 6.304000 | 6.618000 | 3.134000 | 2.448000 |
| std | 9.960637 | 1.734877 | 1.529792 | 1.542996 | 1.858751 | 1.428067 |
| min | 16.000000 | 1.000000 | 2.000000 | 2.000000 | 0.000000 | 0.000000 |
| 25% | 24.000000 | 4.300000 | 5.000000 | 6.000000 | 2.000000 | 1.000000 |
| 50% | 34.000000 | 5.600000 | 6.000000 | 7.000000 | 3.000000 | 2.000000 |
| 75% | 41.000000 | 6.700000 | 7.000000 | 8.000000 | 5.000000 | 3.000000 |
| max | 49.000000 | 10.800000 | 10.000000 | 10.000000 | 9.000000 | 7.000000 |

*Checking missing Data: Checked for null values.

*No missing values detected

```
: <bound method DataFrame.sum of      User_ID    Age   Gender  Daily_Screen_Time(hrs)
 0    False  False  False           False        False
 1    False  False  False           False        False
 2    False  False  False           False        False
 3    False  False  False           False        False
 4    False  False  False           False        False
 ..
 ..
 ..
 495   False  False  False           False        False
 496   False  False  False           False        False
 497   False  False  False           False        False
 498   False  False  False           False        False
 499   False  False  False           False        False
```

*Head (5) for printing initial first 5 columns

| | User_ID | Age | Gender | Daily_Screen_Time(hrs) | Sleep_Quality(1-10) | Stress_Level(1-10) | Days_Without_Social_Media |
|---|---------|-----|--------|------------------------|---------------------|--------------------|---------------------------|
| 0 | U001 | 44 | Male | 3.1 | 7.0 | 6.0 | 2.0 |
| 1 | U002 | 30 | Other | 5.1 | 7.0 | 8.0 | 5.0 |
| 2 | U003 | 23 | Other | 7.4 | 6.0 | 7.0 | 1.0 |
| 3 | U004 | 36 | Female | 5.7 | 7.0 | 8.0 | 1.0 |
| 4 | U005 | 34 | Female | 7.0 | 4.0 | 7.0 | 5.0 |

4. Predicting Happiness Using Linear Regression

- Importing required features for predicting Linear regression.
- Reading adidas CSV file and assign X and Y variables.
- Instantiates a Standard Scaler to standardize the features.
- Scales X and splits it into train/test sets (80 % train, 20 % test) .
- Random State = 42.

```
x = numeric_features.drop('Happiness_Index(1-10)', axis=1)
y = numeric_features['Happiness_Index(1-10)']

x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.3, random_state=42)
```

- Predicts on the test set (y pred).
- Computes and prints the Mean Squared Error and R² score.
- Prints the regression coefficients.

```
y_pred = model.predict(x_test)
r2 = r2_score(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))

print("\nModel Performance:")
print("R2 score:", r2)
```

```
\Model Coefficients: [ 0.00622924 -0.07675173  0.34520363
Model Intercept: 9.33874689962726
```

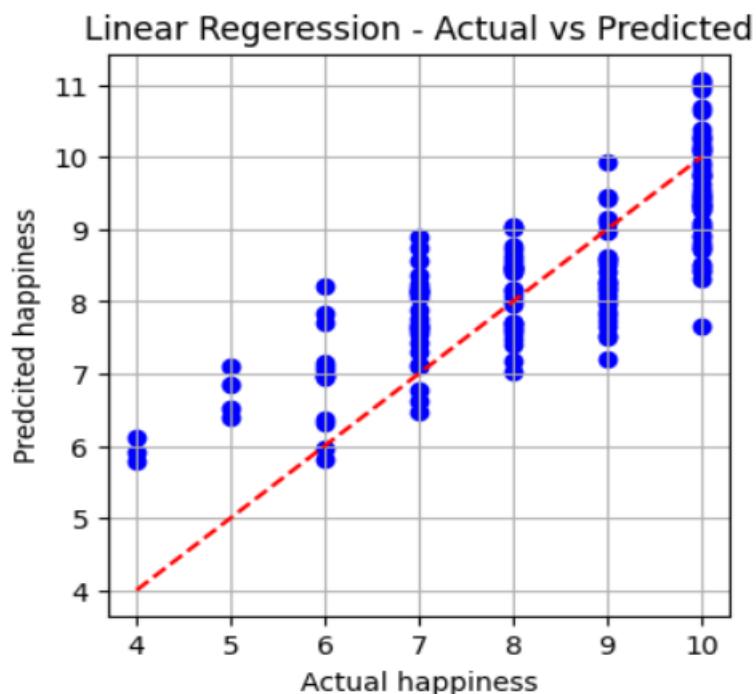
```
Model Performance:
R2 score: 0.9579548353228006
```

- * The model has excellent predictive performance with an R^2 of 0.96
- * Meaning it fits the data very well.
- * The intercept (9.3387) is the baseline prediction features are zero

- Comparing the Actual VS Predicted.

| Actual vs Predicted: | | |
|----------------------|--------|-----------|
| | Actual | Predicted |
| 361 | 7.0 | 8.110613 |
| 73 | 8.0 | 8.471536 |
| 374 | 9.0 | 7.869610 |
| 155 | 10.0 | 9.273720 |
| 104 | 9.0 | 8.004732 |
| .. | ... | ... |
| 266 | 10.0 | 10.392033 |
| 23 | 7.0 | 7.095210 |
| 222 | 8.0 | 9.051211 |
| 261 | 10.0 | 9.464902 |
| 426 | 8.0 | 7.706504 |

```
plt.figure(figsize=(4,4))
plt.scatter(y_test, y_pred, color="blue")
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color="red", linestyle="--")
plt.xlabel("Actual happiness")
plt.ylabel("Predicted happiness")
plt.title("Linear Regression - Actual vs Predicted")
plt.grid(True)
plt.show()
```



*This plot shows a linear regression comparison.
 * Between actual happiness and predicted happiness.
 *Blue dot represent individual observations.
 *The model generally captures the positive trend.
 *The spread of points suggests moderate residual variance.
 *The model isn't perfect but gives a reasonable approximation.

4. Predicting Stress Level Using Lasso Regression

- Importing required features for predicting Lasso regression.
- Reading CSV file and assign X and Y variables.
- Instantiates a Standard Scaler to standardize the features.
- Scales X and splits it into train/test sets (80 % train, 20 % test).
- Random State = 42.

```

df = pd.read_csv("C:/Users/appuv/Desktop/PROJECT/Mental_Health_and_Social_Media_Balance_Dataset.csv")
numeric_features = df.select_dtypes(include=[np.number])
x = numeric_features.drop('Stress_Level(1-10)', axis=1)
y = numeric_features['Stress_Level(1-10)']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.fit_transform(x_test)

```

- Create 50 values of alpha logarithmically.
- LassoCV performs with built-in cross-validation.
- The optimal alpha is retrieved from the alpha-attribute.
- A new lasso model is instantiated with the best alpha.

```

alphas = np.logspace(-3, 1, 50)
lasso_cv = LassoCV(alphas=alphas, cv=5, max_iter=10000)
lasso_cv.fit(x_train, y_train)

best_alpha = lasso_cv.alpha_
print(f"Best alpha from Cv: {best_alpha:.4f}")

lasso = Lasso(alpha=best_alpha)
lasso.fit(x_train, y_train)

```

- Print R2 and MSE

```
Best alpha from Cv: 0.0026
MSE: 0.9044380661404757
R2 score: 0.6244651776530161
```

***Best Alpha CV is 0.0026 the optimal regularization parameter is selected.**

***The Mean Squared Error measures the average squared difference between Predicted and actual values. Lower MSE indicates better fit.**

***The R² is 0.62 indicates how much variance in the dependent variable is explained.**

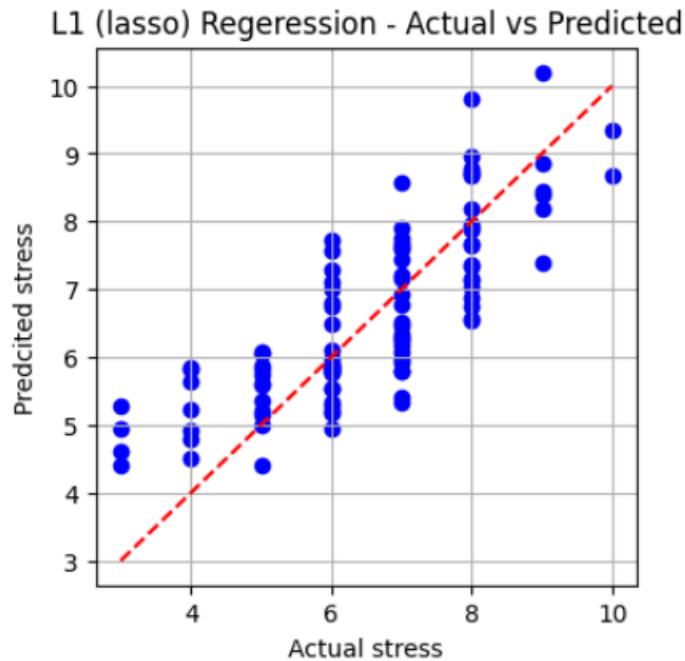
- Comparing the Actual VS Predicted.

| Actual vs Predicted: | | |
|----------------------|--------|-----------|
| | Actual | Predicted |
| 361 | 7.0 | 7.673511 |
| 73 | 6.0 | 6.490530 |
| 374 | 7.0 | 6.507964 |
| 155 | 5.0 | 5.876924 |
| 104 | 6.0 | 6.756837 |
| .. | ... | ... |
| 347 | 8.0 | 7.957511 |
| 86 | 4.0 | 4.513907 |
| 75 | 5.0 | 5.211744 |
| 438 | 6.0 | 7.016593 |
| 15 | 6.0 | 5.852412 |

***The predicted values are continuous decimals, suggesting a regression or probabilistic model.**

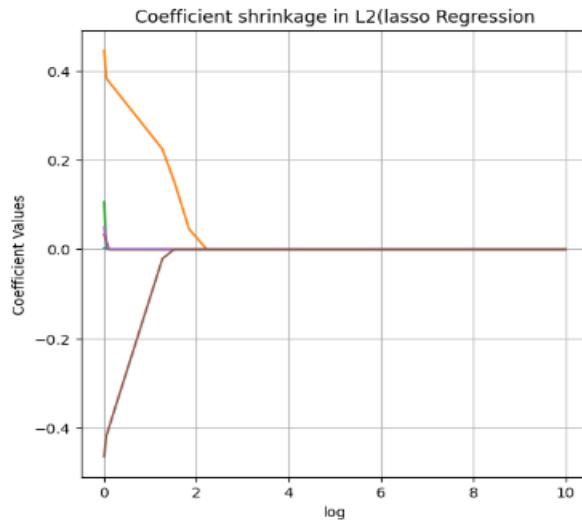
***Difference between actual and predicted indicate the model's error for each instance.**

Plotting the Actual VS Predicted



- *This plot shows a lasso regression comparison.
- * Between actual stress and predicted stress
- *Blue dot represent individual observations.
- *The model generally captures the positive trend.
- *The spread of points suggests moderate residual variance.

- Coefficient Shrinkage



- *X-axis represent the log of the regularization parameter.
- *As log increases the penalty on the coefficient gets stronger
- *Y-axis shows how the regression coefficients change with log.
- *The plot helps choose an optional log by balance bias-variance
- *In ridge regression, coefficients are shrunk but not set exactly to zero unlike lasso which performs variable selection.

5. CONCLUSION

- The Linear regression predicts Happiness accuracy of 96%.
- Lasso Regression gives an accuracy for Stress Level of 62%.
- Happiness is heavily driven by Sleep Quality and Stress Level.